

Transformer and Large Language Models

Instructor: Lei Wu¹

Mathematical Introduction to Machine Learning

Peking University, Fall 2025

¹School of Mathematical Sciences; Center for Machine Learning Research

Outline

- ① Transformer Architecture
- ② Large Language Models (LLMs)
- ③ Vision Transformers (ViT)

Outline

① Transformer Architecture

② Large Language Models (LLMs)

③ Vision Transformers (ViT)

Transformer

Transformers

- were introduced in *Attention is all you need* (Vaswani et al., NeurIPS 2017);
- have revolutionized NLP, CV, robotics and many applications;
- have enabled the creation of powerful LLMs such as GPT-4;
- hold the promise of unlocking the potential for AGI (artificial general intelligence).



Sequence Modeling

Consider a simple block for sequence modeling:

$$(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n) \xrightarrow{\mathcal{T}} (\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_n).$$

In practical models, we may compose the \mathcal{T} -type blocks for many times (aka layers).

Sequence Modeling

Consider a simple block for sequence modeling:

$$(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n) \xrightarrow{\mathcal{T}} (\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_n).$$

In practical models, we may compose the \mathcal{T} -type blocks for many times (aka layers).

- Recurrence

$$\mathbf{y}_i = f(\mathbf{x}_i, \mathbf{y}_{i-1}).$$

Sequence Modeling

Consider a simple block for sequence modeling:

$$(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n) \xrightarrow{\mathcal{T}} (\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_n).$$

In practical models, we may compose the \mathcal{T} -type blocks for many times (aka layers).

- Recurrence

$$\mathbf{y}_i = f(\mathbf{x}_i, \mathbf{y}_{i-1}).$$

- Convolution (local connection)

$$\mathbf{y}_i = f(\mathbf{x}_{i-1}, \mathbf{x}_i, \mathbf{x}_{i+1}).$$

Sequence Modeling

Consider a simple block for sequence modeling:

$$(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n) \xrightarrow{\mathcal{T}} (\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_n).$$

In practical models, we may compose the \mathcal{T} -type blocks for many times (aka layers).

- Recurrence

$$\mathbf{y}_i = f(\mathbf{x}_i, \mathbf{y}_{i-1}).$$

- Convolution (local connection)

$$\mathbf{y}_i = f(\mathbf{x}_{i-1}, \mathbf{x}_i, \mathbf{x}_{i+1}).$$

- Attention (**selective** weighted average):

$$\mathbf{y}_i = f \left(\sum_{j=1}^n w_{i,j}(X) \mathbf{x}_j \right),$$

where $W(X) = (w_{i,j}(X)) \in \mathbb{R}^{n \times n}$ satisfies

$$\sum_{j=1}^n w_{i,j}(X) = 1, \quad W_{i,j}(X) \geq 0 \quad \forall i, j \in [n].$$

Attention Mechanism (Cont'd)

We often call $w_{i,j}(X)$'s the **attention score** and we want

the attention scores $(w_{i,1}(X), w_{i,2}(X), \dots, w_{i,n}(X))$ to be **sparse** (i.e., **selective**).

- Attention in vision modeling:



Attention Mechanism (Cont'd)

Attention in machine translation (**cross attention**) ²:

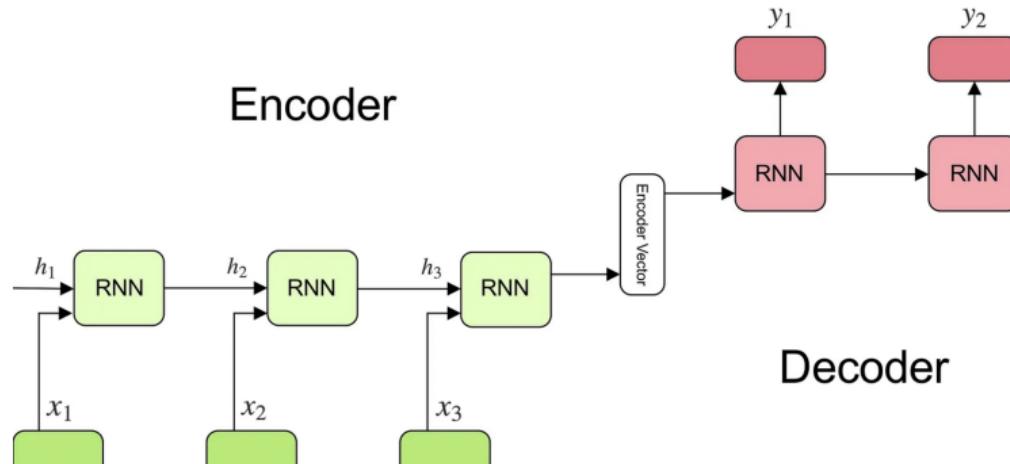


Figure 1: See a better animation in this link.

²Bahdanau et al., Neural Machine Translation by Jointly Learning to Align and Translate, ICLR 2015.

Self-Attention via Dot-Product

- Let $X = (x_1, \dots, x_n) \in \mathbb{R}^{d \times n}$ be our input sequence. We often call $\{x_i\}$ **tokens**.
- A self-attention $\mathbb{A} : \mathbb{R}^{d \times n} \mapsto \mathbb{R}^{n \times n}$ outputs an attention-score map $P = \mathbb{A}(X)$. The most popular choice is

$$\mathbb{A}_{W_K, W_Q}(X) = \sigma \left(\frac{1}{\sqrt{d_{\text{key}}}} (W_K X)^{\top} (W_Q X) \right) \in \mathbb{R}^{n \times n},$$

where

- $W_K, W_Q \in \mathbb{R}^{d_{\text{key}} \times d}$ are the **key** and **query** weight matrices, which are learned from data.
- σ denotes the softmax normalization performed in a column-wise manner, ensuring the column represent a selective average.

The diagram illustrates the computation of the attention matrix P . It starts with a sequence of tokens represented as a vector:

Anthony Hopkins
admired Michael Bay
as a great director.

This vector is multiplied by the query weight matrix W_Q (labeled Q in the diagram) and then by the key weight matrix W_K (labeled K in the diagram). The result is then divided by $\sqrt{d_k}$. The final output is the attention matrix P , which is a matrix where each row represents the attention scores for each token, summing up to 1. An orange arrow points to the last row of the matrix with the text: "Importantly, the sum of each row is 1."

Self-Attention via Dot-Product (Cont'd)

- The dot-products are implemented in a **token-wise manner** (can be naively paralleled):

$$\mathbf{k}_i = W_K \mathbf{x}_i, \quad \mathbf{q}_j = W_Q \mathbf{x}_j \quad \text{for } i, j \in [n]$$
$$(\mathbb{A}_{W_K, W_Q}(X))_{i,j} = \frac{e^{\mathbf{k}_i^\top \mathbf{q}_j}}{\sum_{i'=1}^n e^{\mathbf{k}_{i'}^\top \mathbf{q}_j}}$$

- The attention scores are determined by the **dot-product correlation** among tokens.
- A single-head attention layer $\text{SA} : \mathbb{R}^{d \times n} \mapsto \mathbb{R}^{d \times n}$ is given as follows

$$\text{SA}_{W_K, W_Q, W_V}(X) = V \sigma(K^\top Q),$$

where Q, K, V are called the query, key, value matrices, respectively and given by

$$Q = W_Q X \in \mathbb{R}^{d \times n}, \quad K = W_K X \in \mathbb{R}^{d \times n}, \quad V = W_V X \in \mathbb{R}^{d \times n}.$$

A Transformer Block

- A transformer block defines a sequence-to-sequence map

$$X = (\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n) \in \mathbb{R}^{d \times n} \mapsto Y = (\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_n) \in \mathbb{R}^{d \times n}.$$

- This map consists of two blocks:

$$Y = \mathbb{F}(X + \text{MHA}(X)),$$

where

- **Multi-head attention (MHA)**

$$\text{MHA}(X) := \sum_{h=1}^H W_O^h \text{SA}^h(X).$$

- **Tokenwise feed-forward networks (FFN):**

$$\mathbb{F}(Z) := (h(\mathbf{z}_1), h(\mathbf{z}_2), \dots, h(\mathbf{z}_n)) \in \mathbb{R}^{d \times n}.$$

In practice, $h : \mathbb{R}^d \mapsto \mathbb{R}^d$ is often chosen to be a two-layer MLP with hidden size d_{FF} .

$$h(\mathbf{z}) = W_1^\top \text{ReLU}(W_2 \mathbf{z} + \mathbf{b}),$$

where $W_1, W_2 \in \mathbb{R}^{d_{\text{FF}} \times d}$ and $\mathbf{b} \in \mathbb{R}^d$.

Transformer

- **Input:** Linear embedding to change the dimension of each token.

$$X^{(0)} = VX \text{ with } V \in \mathbb{R}^{d_{\text{model}} \times d}.$$

- **Main block:**

$$X^\ell = \mathbb{F}^{(\ell)}(X^{(\ell-1)} + \text{MHA}^{(\ell)}(X^{(\ell-1)})), \quad 1 \leq \ell \leq L.$$

- **Output:** The output format depends on the tasks. In classification, we may

$$f(X) = p(\mathbf{x}_1^{(L)}),$$

where p can be either a linear layer or small MLP.

- Architecture hyperparameters: d_{model} , H , L , d_{key} , d_{FF} . In practice, a common choice $d_{\text{FF}} = 4d_{\text{model}}$, $d_{\text{key}} = d_{\text{model}}/H$.

Cost Analysis of Transformer

In practice, it is common to choose

$$d_{\text{key}} = d_{\text{model}}/H, \quad d_{\text{FF}} = 4d_{\text{model}}.$$

- **Storage cost** (per Transformer layer):

$$4d_{\text{model}}^2 \text{ (MHA)} + 8d_{\text{model}}^2 \text{ (FFN)}.$$

- **Computation cost:**

- **MHA:**

$$\underbrace{4nd_{\text{model}}^2}_{\text{Q/K/V projections}} + \underbrace{d_{\text{model}} n^2}_{\text{attention scores}}$$

- **FFN:**

$$8d_{\text{model}}^2 n.$$

- **Parallelism:** Tokenwise operations are fully parallelizable.
- **Critical bottleneck:** The total cost grows **quadratically in the sequence length n .** This quadratic attention cost becomes the dominant bottleneck—especially for inference.

Absolute Positional Embedding (APE)

Transformers without positional information are **permutation-equivariant**: reordering the input tokens simply reorders the outputs. To encode order, we need to inject positional information.

The most natural way of injecting position information is using **absolute positional embedding** (APE): let $\mathbf{r}_i \in \mathbb{R}^d$ denote the **position information** for token i :

$$\mathbf{x}_i \rightarrow \mathbf{x}_i + \mathbf{r}_i,$$

- Learnable APE: \mathbf{r}_i are parameters to be learned.
- One-hot APE: $\mathbf{r}_i = \mathbf{e}_i$ where \mathbf{e}_i is the one-hot label with 1 in the i -th coordinates and zero else.
- Sinusoidal APE:

$$\mathbf{r}_i = \left(\sin(i), \cos(i), \sin(i/c), \cos(i/c), \dots, \sin(i/c^{2k/d}), \cos(i/c^{2k/d}) \right) \in \mathbb{R}^d,$$

where c is constant, e.g. 1000.

Absolute Positional Embedding (APE)

Transformers without positional information are **permutation-equivariant**: reordering the input tokens simply reorders the outputs. To encode order, we need to inject positional information.

The most natural way of injecting position information is using **absolute positional embedding** (APE): let $\mathbf{r}_i \in \mathbb{R}^d$ denote the **position information** for token i :

$$\mathbf{x}_i \rightarrow \mathbf{x}_i + \mathbf{r}_i,$$

- Learnable APE: \mathbf{r}_i are parameters to be learned.
- One-hot APE: $\mathbf{r}_i = \mathbf{e}_i$ where \mathbf{e}_i is the one-hot label with 1 in the i -th coordinates and zero else.
- Sinusoidal APE:

$$\mathbf{r}_i = \left(\sin(i), \cos(i), \sin(i/c), \cos(i/c), \dots, \sin(i/c^{2k/d}), \cos(i/c^{2k/d}) \right) \in \mathbb{R}^d,$$

where c is constant, e.g. 1000.

However, APE is rarely used in practice anymore due to:

- APE can not handle input sequence longer than that used in training.
- In many real problems, it is “relative distance” that matters.

Relative Positional Embedding (RPE)

- **Additive RPE:** Let $E = (W_K X)^\top (W_Q X) \in \mathbb{R}^{n \times n}$ be the pre-softmax attention logits. Then, we inject relative position information by

$$\mathbb{A}(X) = \sigma(E - P),$$

where $P = (h(j - i))_{i,j} \in \mathbb{R}^{n \times n}$.

Relative Positional Embedding (RPE)

- **Additive RPE:** Let $E = (W_K X)^\top (W_Q X) \in \mathbb{R}^{n \times n}$ be the pre-softmax attention logits. Then, we inject relative position information by

$$\mathbb{A}(X) = \sigma(E - P),$$

where $P = (h(j-i))_{i,j} \in \mathbb{R}^{n \times n}$.

- **T5 RPE:**

$$h(t) = \begin{cases} |t| & \text{if } |t| \leq B/2 \\ \frac{B}{2} + \frac{B}{2} \left\lfloor \frac{\log(\frac{|t|}{B/2})}{\log(\frac{D}{B/2})} \right\rfloor & \text{if } \frac{B}{2} \leq |t| \leq D \\ B-1 & \text{if } |t| \geq D \end{cases}$$

where where B is the total number of relative position buckets and D is the maximum distance we explicitly bucketize.

- **Alibi RPE:** $h(t) = -\alpha|t| + \beta$, where the α and β can be either learnable or fixed.

Relative Positional Embedding (RPE)

- **Additive RPE:** Let $E = (W_K X)^\top (W_Q X) \in \mathbb{R}^{n \times n}$ be the pre-softmax attention logits. Then, we inject relative position information by

$$\mathbb{A}(X) = \sigma(E - P),$$

where $P = (h(j - i))_{i,j} \in \mathbb{R}^{n \times n}$.

- **T5 RPE:**

$$h(t) = \begin{cases} |t| & \text{if } |t| \leq B/2 \\ \frac{B}{2} + \frac{B}{2} \left\lfloor \frac{\log(\frac{|t|}{B/2})}{\log(\frac{D}{B/2})} \right\rfloor & \text{if } \frac{B}{2} \leq |t| \leq D \\ B - 1 & \text{if } |t| \geq D \end{cases}$$

where where B is the total number of relative position buckets and D is the maximum distance we explicitly bucketize.

- **Alibi RPE:** $h(t) = -\alpha|t| + \beta$, where the α and β can be either learnable or fixed.
- Currently, the most popular RPE is the rotary positional embedding (**RoPE**), which has been adopted in nearly all large language models.

Readings

- The original paper <https://arxiv.org/abs/1706.03762>
- Annotated Transformer <https://jalammar.github.io/illustrated-transformer/>
- Illustrated Transformer <https://poloclub.github.io/transformer-explainer/>

Outline

- 1 Transformer Architecture
- 2 Large Language Models (LLMs)
- 3 Vision Transformers (ViT)

Tokenization

Before applying embeddings, natural language text must first be **tokenized**.

- **Definition:** Tokenization is the process of converting raw text into a sequence of discrete units called *t*okens.
- **Purpose:** It transforms unstructured text into a form that models can process, enabling downstream operations such as embedding, attention, and sequence modeling.

Tokenization

Before applying embeddings, natural language text must first be **tokenized**.

- **Definition:** Tokenization is the process of converting raw text into a sequence of discrete units called *tokens*.
- **Purpose:** It transforms unstructured text into a form that models can process, enabling downstream operations such as embedding, attention, and sequence modeling.

Types of Tokenization

- **Word-level Tokenization:** Splits text into individual words and punctuation.

"Transformers are amazing!" → ["Transformers", "are", "amazing", "!"]

Tokenization

Before applying embeddings, natural language text must first be **tokenized**.

- **Definition:** Tokenization is the process of converting raw text into a sequence of discrete units called *tokens*.
- **Purpose:** It transforms unstructured text into a form that models can process, enabling downstream operations such as embedding, attention, and sequence modeling.

Types of Tokenization

- **Word-level Tokenization:** Splits text into individual words and punctuation.
"Transformers are amazing!" -> ["Transformers", "are", "amazing", "!"]
- **Subword-level Tokenization:** Decomposes words into smaller units (subwords), allowing robust **handling of rare words, morphology, and open vocabularies**.

"Transformers are amazing!" ->
["Trans", "#form", "#ers", "are", "amaz", "#ing", "!"]

Tokenization

Before applying embeddings, natural language text must first be **tokenized**.

- **Definition:** Tokenization is the process of converting raw text into a sequence of discrete units called *tokens*.
- **Purpose:** It transforms unstructured text into a form that models can process, enabling downstream operations such as embedding, attention, and sequence modeling.

Types of Tokenization

- **Word-level Tokenization:** Splits text into individual words and punctuation.
"Transformers are amazing!" -> ["Transformers", "are", "amazing", "!"]
- **Subword-level Tokenization:** Decomposes words into smaller units (subwords), allowing robust **handling of rare words, morphology, and open vocabularies**.
"Transformers are amazing!" ->
["Trans", "#form", "#ers", "are", "amaz", "#ing", "!"]
- **In practice:** Modern LLMs rely on standardized subword tokenizers, for example those implemented in Hugging Face's transformers library.

Training LLMs Requires Many Stabilization Tricks

- **Scaled dot-product attention** Scaling by $1/\sqrt{d_k}$ controls the variance of attention scores:

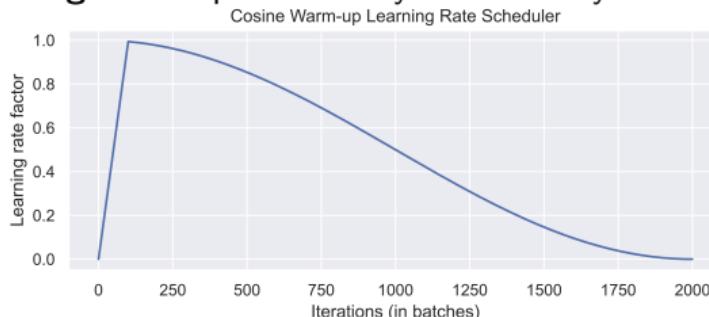
$$\mathbb{A}_{W_K, W_Q}(X) = \sigma\left(\frac{1}{\sqrt{d_k}}(W_K X)^\top (W_Q X)\right) \in \mathbb{R}^{n \times n}.$$

- **Normalization** (LayerNorm or RMSNorm) **with residual connections**: These are essential for stable signal propagation across depth.

$$\tilde{X}^{(\ell-1)} = \text{Norm}\left(X^{(\ell-1)}\right),$$

$$X^{(\ell)} = \tilde{X}^{(\ell-1)} + \mathbb{F}\left(\text{MHA}(\tilde{X}^{(\ell-1)})\right).$$

- **AdamW** (Adam+decoupled weight decay) optimizer with $\beta_1 = 0.9$ and $\beta_2 \in \{0.95, 0.98\}$, together with **gradient clipping**.
- **Learning-rate scheduling**: warmup followed by cosine decay.



Bidirectional Encoder Representations from Transformers (BERT)

- Developed by Google (Devlin et al., 2018).
- **Bidirectional encoder:** Unlike traditional left-to-right or right-to-left language models, BERT uses *deep bidirectional attention*, allowing it to incorporate context from both sides simultaneously.

Bidirectional Encoder Representations from Transformers (BERT)

- Developed by Google (Devlin et al., 2018).
- **Bidirectional encoder:** Unlike traditional left-to-right or right-to-left language models, BERT uses *deep bidirectional attention*, allowing it to incorporate context from both sides simultaneously.

Pre-training Tasks

- Masked Language Modeling (MLM): Randomly masks a subset of tokens and predicts the missing ones.
- Next Sentence Prediction (NSP): Predicts whether two sentences appear consecutively in the original text.

(Later work shows that NSP provides little benefit and can often be removed.)

Bidirectional Encoder Representations from Transformers (BERT)

- Developed by Google (Devlin et al., 2018).
- **Bidirectional encoder:** Unlike traditional left-to-right or right-to-left language models, BERT uses *deep bidirectional attention*, allowing it to incorporate context from both sides simultaneously.

Pre-training Tasks

- Masked Language Modeling (MLM): Randomly masks a subset of tokens and predicts the missing ones.
- Next Sentence Prediction (NSP): Predicts whether two sentences appear consecutively in the original text.

(Later work shows that NSP provides little benefit and can often be removed.)

Fine-tuning

- The pre-trained encoder can be adapted to downstream tasks such as question answering, sentiment analysis, text classification, etc.

Generative Pre-trained Transformer (GPT)

- **Next-token prediction** (autoregressive model):

$$\max_{\theta} \frac{1}{n} \sum_{i=1}^n \log P_{\theta}(x_i | x_1, \dots, x_{i-1}).$$

Generative Pre-trained Transformer (GPT)

- **Next-token prediction** (autoregressive model):

$$\max_{\theta} \frac{1}{n} \sum_{i=1}^n \log P_{\theta}(x_i | x_1, \dots, x_{i-1}).$$

- Text Generation:

```
text = [ $\langle bos \rangle$ ] or [some context]
```

```
while True:
```

```
    logit = decoder(embed(text))
```

```
    index = top(logit[-1])
```

```
    token = vocabulary(index)
```

```
    if token ==  $\langle eos \rangle$  :
```

```
        break
```

```
    text.append(token)
```

```
return text
```

Emergent Abilities: In-Context Learning and Prompting

- Traditional pre-train → fine-tune pipelines require updating model parameters.
- GPT's autoregressive training objective leads to **emergent abilities** that BERT-style models largely lack.
- **In-Context Learning (ICL)**: The model can learn from examples provided in the input.

In the following lines, the symbol \rightarrow represents a simple mathematical operation.

$100 + 200 \rightarrow 301$

$838 + 520 \rightarrow 1359$

$343 + 128 \rightarrow 472$

$647 + 471 \rightarrow 1119$

$64 + 138 \rightarrow 203$

$498 + 592 \rightarrow$

Answer:

1091

- **Prompting**: Solve tasks by giving instructions or demonstrations without parameter updates.

Next-Token Prediction Training Yields Emergent Abilities

- **Training objective:** GPT is trained to model

$$P(x_i \mid x_1, x_2, \dots, x_{i-1})$$

which requires predicting the next token conditioned on *any* prefix.

Next-Token Prediction Training Yields Emergent Abilities

- **Training objective:** GPT is trained to model

$$P(x_i \mid x_1, x_2, \dots, x_{i-1})$$

which requires predicting the next token conditioned on *any* prefix.

- **Prefixes = implicit tasks:** During training, prefixes may contain:
 - instructions,
 - demonstrations,
 - question–answer pairs,
 - multi-step reasoning patterns.

The model must interpret these patterns to predict the next token correctly.

Next-Token Prediction Training Yields Emergent Abilities

- **Training objective:** GPT is trained to model

$$P(x_i \mid x_1, x_2, \dots, x_{i-1})$$

which requires predicting the next token conditioned on *any* prefix.

- **Prefixes = implicit tasks:** During training, prefixes may contain:

- instructions,
- demonstrations,
- question–answer pairs,
- multi-step reasoning patterns.

The model must interpret these patterns to predict the next token correctly.

- **Task inference emerges:** To minimize training loss, the model implicitly learns to:

- infer the task from the prefix,
- adapt behavior based on examples in the input,
- follow natural-language instructions.

Next-Token Prediction Training Yields Emergent Abilities

- **Training objective:** GPT is trained to model

$$P(x_i \mid x_1, x_2, \dots, x_{i-1})$$

which requires predicting the next token conditioned on *any* prefix.

- **Prefixes = implicit tasks:** During training, prefixes may contain:

- instructions,
- demonstrations,
- question–answer pairs,
- multi-step reasoning patterns.

The model must interpret these patterns to predict the next token correctly.

- **Task inference emerges:** To minimize training loss, the model implicitly learns to:

- infer the task from the prefix,
- adapt behavior based on examples in the input,
- follow natural-language instructions.

- **Resulting abilities:**

- **In-Context Learning (ICL):** learning from examples in the prompt;
- **Prompting:** solving diverse tasks without parameter updates;
- **Generalization beyond language:** symbolic reasoning, coding, planning, etc.

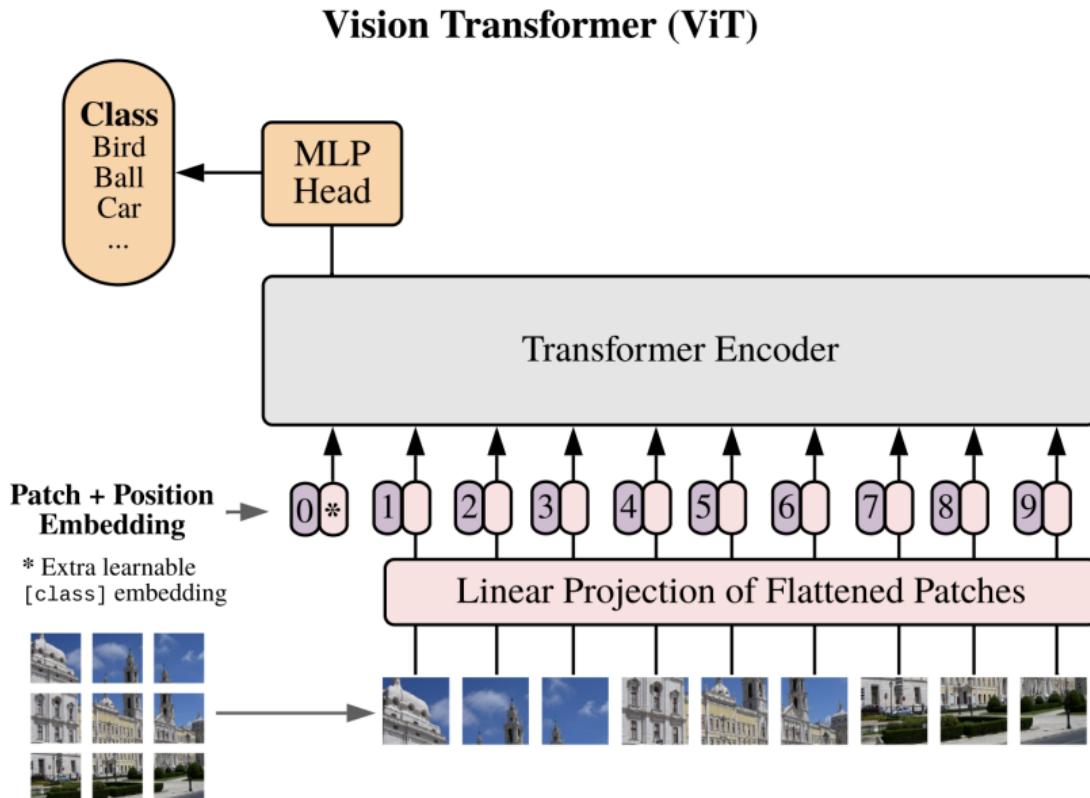
Remark

GPT and its focus on next-token prediction have fundamentally transformed how pre-trained models are utilized, marking a significant step toward AGI. The transition from BERT to GPT represents a **major breakthrough** in this evolution.

Outline

- 1 Transformer Architecture
- 2 Large Language Models (LLMs)
- 3 Vision Transformers (ViT)

Vision Transformer (ViT)



Summary

- **Transformer Architecture** Self-attention + MHA + FFN form a fully parallel sequence model; positional encodings inject order information.

Summary

- **Transformer Architecture** Self-attention + MHA + FFN form a fully parallel sequence model; positional encodings inject order information.
- **Large Language Models** GPT-style autoregressive training + modern optimization lead to in-context learning and prompting.

Summary

- **Transformer Architecture** Self-attention + MHA + FFN form a fully parallel sequence model; positional encodings inject order information.
- **Large Language Models** GPT-style autoregressive training + modern optimization lead to in-context learning and prompting.
- **Vision Transformer (ViT)** Applying Transformer to image patches shows the architecture's universality.