Theoretical Deep Learning

# Lecture 4: Kernel methods, representer theorem and RKHSs

July 29, 2021

*Lecturer: Lei Wu* *Scribe: Lei Wu*

## 1 Feature-based methods

We start by considering the linear regression, for which the hypothesis class is

$$\mathcal{F} = \{\beta^T x : \beta \in \mathbb{R}^d\},$$

where we omit the bias term for simplicity. The ridge regression penalizes the squared $\ell_2$ norm of $\beta$:

$$\hat{\beta}_n = \operatorname*{argmin}_{\beta\in\mathbb{R}^d} \frac{1}{n}\sum_{i=1}^n (\beta^T x_i - y_i)^2 + \lambda\|\beta\|_2^2.$$

The minimizer has a closed-form solution:

$$\hat{\beta}_n = \left(\frac{1}{n}X^T X + I\right)^{-1}\frac{1}{n}Xy,$$

where $X = (x_1, \ldots, x_n)^T \in \mathbb{R}^{n\times d}$, $y = (y_1, \ldots, y_n)^T \in \mathbb{R}^n$. Another population one is LASSO, which penalizes the $\ell_1$ norm of parameters:

$$\min_{\beta\in\mathbb{R}^d} \frac{1}{n}\sum_{i=1}^n (\beta^T x_i - y_i)^2 + \lambda\|\beta\|_1.$$

To consider nonlinear functions, we can consider the model:

$$f(x;\beta) = \sum_{j=1}^m \beta_j \varphi_j(x).$$

Here, $\varphi_1, \ldots, \varphi_n$ are a set of (nonlinear) basis functions, which are often referred to as *features* in machine learning. Accordingly, the feature map is defined as $\Phi : \mathcal{X} \mapsto \mathbb{R}^m$ with $\Phi(x) = (\varphi_1(x), \ldots, \varphi_n(x))^T \in \mathbb{R}^m$. Typical examples includes

- Spectral methods: $\{\varphi_j\}$ are either Fourier basis or orthogonal polynomials.
- Splines: $\{\varphi_j\}$ are piecewise polynomials.
- Computer vision: Some hand-crafted features.

We can consider two types of feature-based methods.

$$\frac{1}{n}\sum_{i=1}^n (\beta^T \Phi(x_i) - y_i)^2 + \lambda\|\beta\|_2^2$$
$$\frac{1}{n}\sum_{i=1}^n (\beta^T \Phi(x_i) - y_i)^2 + \lambda\|\beta\|_1.$$

## 1.1 General feature-based methods

**$\ell_2$ extension.** The previous idea can be extended to a general feature-based model:

$$f(x;\beta) = \langle \beta, \Phi(x) \rangle_{\mathcal{H}}, \tag{1.1}$$

where

- $\mathcal{H}$ is the feature space, which can be any Hilbert space;
- $\Phi : \mathcal{X} \mapsto \mathcal{H}$ is the feature map;
- The "coefficients" are $\beta \in \mathcal{H}$.

If taking $\mathcal{H} = \mathbb{R}^m$ and $\Phi(x) = (\phi(x_1), \dots, \phi(x_n))^T \in \mathbb{R}^m$, we recover the classical ones. However, the advantage of the formulation (1.1) is that it includes the case where $m = \infty$. Below is an example:

**Random feature models (RFMs).** Consider

$$f(x;\beta) = \int \beta(w)\varphi(x;w)\,\mathrm{d}\pi(w) = \langle \beta, \varphi(x;\cdot) \rangle_{L^2(\pi)}, \tag{1.2}$$

where $\pi$ is a fixed distribution. In this case, the feature map is given by

$$\Phi : \mathcal{X} \mapsto L^2(\pi), \qquad \Phi(x) = \varphi(x;\cdot),$$

and the parameter is $\beta \in L^2(\pi)$. The model (1.2) can be viewed as the continuum limit of the following random feature model

$$f(x;\beta) = \frac{1}{m}\sum_{j=1}^{m} \beta_j \varphi(x;w_j),$$

where $w_1, \dots, w_m$ are independently sampled from $\pi$ and fixed.

Now, the model (1.1) is well-defined. The objective function of the corresponding ridge regression can be written as

$$\hat{\mathcal{R}}_n(\beta) = \frac{1}{n}\sum_{i=1}^{n} (\langle \beta, \Phi(x_i) \rangle_{\mathcal{H}} - y_i)^2 + \lambda \|\beta\|_{\mathcal{H}}^2. \tag{1.3}$$

How can we optimize (1.3), which is an infinitely dimensional problem?

## 1.2 $\ell_1$ extension.

Consider the random feature methods:

$$\frac{1}{n}\sum_{i=1}^{n} \left( \sum_{j=1}^{m} \beta_j \varphi(x_i;w_j) - y_i \right)^2 + \frac{\lambda}{m}\sum_{j=1}^{m} |\beta_j|.$$

Assume for any $x \in \mathcal{X}$, $\operatorname{ess\,sup}_w |\varphi(x;w)| < \infty$. Then, the continuum limit of the above method is given by

$$\min_{\beta \in L^1(\pi)} \frac{1}{n}\sum_{i=1}^{n} \left( \int \beta(w)\varphi(x_i;w)\,\mathrm{d}\pi(w) - y_i \right)^2 + \lambda \int |\beta(w)|\,\mathrm{d}\pi(w).$$

This method can not be analyzed using the kernel theory.

## 2 Representer theorem and kernel methods

When it is clear from the context, we will drop the subscripts in $\langle , \rangle_{\mathcal{H}}$ and $\| \cdot \|_{\mathcal{H}}$ for simplicity. Let us consider a general problem:

$$\hat{R}_n(\beta) = \frac{1}{2n} \sum_{i=1}^{n} \ell(f(x_i; \beta), y_i) + \lambda r(\|\beta\|), \tag{2.1}$$

- $f(x; \beta) = \langle \beta, \Phi(x) \rangle$
- $\ell$ is a general loss function.
- $r : [0, \infty) \mapsto [0, \infty)$ is a strictly increasing function.

**Theorem 2.1** (Representer theorem)**.** *Let $\hat{\beta}$ the a minimizer of (2.1). Then, there must exist $a_1, \ldots, a_n \in \mathbb{R}$ such that $\hat{\beta} = \sum_{i=1}^{n} a_i \Phi(x_i)$ and*

$$f(x; \hat{\beta}) = \langle \hat{\beta}, \Phi(x) \rangle = \sum_{i=1}^{n} a_i k(x_i, x), \tag{2.2}$$

*where $k(x, x') := \langle \Phi(x), \Phi(x') \rangle$.*

*Proof.* Let $V_n = \text{span}\{\Phi(x_1), \ldots, \Phi(x_n)\} \subset \mathcal{H}$. For any $\beta \in \mathcal{H}$, we can decompose it as follows

$$\beta = \beta_{\|} + \beta_{\perp},$$

where $\beta_{\|} \in V_n, \beta_{\perp} \in V_n^{\perp}$. Hence, $\|\beta\|^2 = \|\beta_{\|}\|^2 + \|\beta_{\perp}\|^2$. Since $r(\cdot)$ is non-decreasing,

$$r(\|\beta\|) \geq r(\|\beta_{\|}\|). \tag{2.3}$$

On the other hand, for any $x_i$,

$$f(x_i; \beta) = \langle \beta, \Phi(x_i) \rangle = \langle \beta_{\|}, \Phi(x_i) \rangle + \langle \beta_{\perp}, \Phi(x_i) \rangle = \langle \beta_{\|}, \Phi(x_i) \rangle, \tag{2.4}$$

where the last equality is due to $\beta_{\perp} \in V_n^{\perp}$. Combining (2.3) and (2.4), we have for any $\beta \in \mathcal{H}$,

$$\hat{\mathcal{R}}_n(\beta) \geq \hat{\mathcal{R}}_n(\beta_{\|}).$$

Therefore, we can take $\hat{\beta}_{\|} = \sum_{i=1}^{n} a_i \Phi(x_i)$. Then, the function represented can be written as

$$f(x; \beta) = \langle \hat{\beta}_{\|}, \Phi(x) \rangle = \sum_{i=1}^{n} a_i \langle \Phi(x_i), \Phi(x) \rangle = \sum_{i=1}^{n} a_i k(x_i, x).$$

□

This theorem allows transforming the infinite-dimensional optimization problem (2.1) into a finite dimensional problem. Moreover, we only need to access the kernel $k(\cdot, \cdot)$ without needing to evaluate the feature maps.

**The reduced model.** Representer theorem implies that we only need to choose

$$\beta = \sum_{j=1}^{n} a_j \Phi(x_j), \qquad f(x;\beta) = \sum_{j=1}^{n} a_j k(x_j, x).$$

Moreover,

$$\|\beta\|^2 = \langle \sum_{j=1}^{n} a_j \Phi(x_j), \sum_{j=1}^{n} a_j \Phi(x_j) \rangle = \sum_{i,j=1}^{n} k(x_i, x_j) a_i a_j = a^T K a,$$

where $a = (a_1, \ldots, a_n)^T \in \mathbb{R}^n$ and $K = (k(x_i, x_j)) \in \mathbb{R}^{n \times n}$ is the *kernel matrix*.

The kernel ridge regression (KRR) corresponds to the case where $\ell(y, y') = (y - y')^2$ and $r(t) = t^2$, i.e., the problem (1.3). Then, the problem can be reduced to the following $n$-dimensional problem

$$\hat{\mathcal{R}}_n(a) = \frac{1}{n} \|Ka - y\|_2^2 + \lambda a^T K a, \tag{2.5}$$

whose solution is given by

$$a = (\frac{1}{n} K + I)^{-1} y.$$

In general, kernel methods refer to methods whose hypothesis class is given by

$$\mathcal{F} = \left\{ \sum_{j=1}^{n} a_j k(x_j, \cdot) : a \in \mathbb{R}^n \right\}.$$

Mathematically, the kernel is defined as follows.

**Definition 2.2** (kernel)**.** $k : \mathcal{X} \times \mathcal{X} \mapsto \mathbb{R}$ is said to be a kernel if there exists a feature map $\Phi : \mathcal{X} \mapsto \mathcal{H}$ such that

$$k(x, x') = \langle \Phi(x), \Phi(x') \rangle.$$

Below is a list of popular kernels.

**Polynomial kernel:** $k(x, x') = (1 + x^T x')^s$ is a kernel for any $s \in \mathbb{N}_+$.

- Linear ($s = 1$). We have $k(x, x') = \langle \Phi(x), \Phi(x') \rangle$ with

$$\Phi(x) = (1, x_1, \ldots, x_d).$$

- Quadratic ($s = 2$): The feature map is given by

$$\Phi(x) = (\underbrace{x_d^2, \ldots, x_1^2}_{\text{quadratic}}, \underbrace{\sqrt{2}x_d x_{d-1}, \ldots, \sqrt{2}x_d x_1, \sqrt{2}x_{d-1}x_{d-2}, \ldots, \sqrt{2}x_2 x_1}_{\text{cross terms}}, \underbrace{\sqrt{2}x_d, \ldots, \sqrt{2}x_1}_{\text{linear terms}}, \underbrace{1}_{\text{constant}}).$$

$$\begin{aligned} \langle \Phi(x), \Phi(x') \rangle &= \sum_{i=1}^{d} (x_i)^2 (x_i')^2 + 2 \sum_{i \neq j} x_i x_j x_i' x_j' + 2 \sum_i x_i x_i' + 1 \\ &= (\sum_{i=1}^{d} x_i x_i')^2 + 2 \sum_i x_i x_i' + 1 \\ &= (x^T x' + 1)^2 \end{aligned} \tag{2.6}$$

**Gaussian kernel:** $k(x, x') = e^{-\frac{\|x-x'\|_2^2}{2}}$. Considering $d = 1$, we have

$$k(x, x') = e^{-\frac{x^2}{2}-\frac{x'^2}{2}} e^{xx'} = e^{-\frac{x^2}{2}-\frac{x'^2}{2}} \sum_n \frac{1}{n!}(x)^n (x')^n$$
$$= \langle \Phi(x), \Phi(x) \rangle,$$

where $\Phi(x) = e^{-\frac{x^2}{2}}(1, x, \frac{1}{\sqrt{2}}x^2, \ldots, \frac{1}{\sqrt{n!}}x^n, \ldots)$.

**Laplace kernel:**

$$k(x, x') = e^{-\frac{\|x-x'\|_2}{\sigma}}.$$

This kernel is less smooth than the Gaussian kernel. Recently, it has been shown that the Laplace kernel is intimately related to neural network models in the kernel regime.

For a specific problem, choosing appropriate kernels is highly non-trivial. One may need to incorporate the domain knowledge into the kernel design.

# 3 Reproducing kernel Hilbert spaces

In this section, we ask the question:

*What kind of functions can be "efficiently" learned by kernel methods?*

By representer theorem, consider

$$\mathcal{F} = \cup_{n=1}^{\infty} \mathcal{F}_n,$$

where

$$\mathcal{F}_n = \left\{ \sum_{j=1}^{n} a_j k(\cdot, x_j) \,:\, x_j \in \mathcal{X}, a_j \in \mathbb{R}, j \in [n] \right\}.$$

This intuition tells us that what kind of functions can be "approximated" by kernel methods. We are interested in functions $f \in \bar{\mathcal{F}}$. However, the problem is how to take the closure and measure the complexity of $f \in \bar{\mathcal{F}}$? Without imposing constraints on the norm of coefficients $\{a_j\}$ in taking the closure, this space can be extremely large. For example, if the corresponding features are polynomials, then $\mathcal{F}$ contains all the continuous functions because of the Stone-Weierstrass theorem. However, $C(\mathcal{X})$ is too large since the Rademacher complexity is $O(1)$. We hope that the Rademacher complexity is on the order of $O(1/\sqrt{n})$.

We need to define an "appropriate" norm for $f \in \mathcal{F}$.

Let us take a step back to the feature-based representation:

$$\beta = \sum_{j=1}^{n} a_j \Phi(x_j), \qquad f(x; \beta) = \sum_{j=1}^{n} a_j k(x_j, \cdot).$$

In KRR, we penalize $\|\beta\|^2$ of the hypothesis. This means that $\|\beta\|^2$ should be a good norm of the represented function $f(x; \beta)$, i.e,

$$\|f(\cdot; \beta)\|^2 = \|\beta\|^2 = \langle \sum_{i=1}^{n} a_i \Phi(x_i), \sum_{j=1}^{n} a_j \Phi(x_j) \rangle = \sum_{i,j=1}^{n} a_i a_j k(x_i, x_j).$$

This intuition can be made rigorous by the following theorem.

**Theorem 3.1** (Moore-Aronsajn theorem)**.** *Let $k : \mathcal{X} \times \mathcal{X} \mapsto \mathbb{R}$ be any kernel. Let $\mathcal{H}^0 = span(\{k(\cdot, x) : x \in \mathcal{X}\})$ and endow it with the inner product:*

$$\langle f, g\rangle_{\mathcal{H}^0} = \sum_{i=1}^{n}\sum_{j=1}^{m} \alpha_i\beta_j k(x_i, x'_j), \tag{3.1}$$

*where $f = \sum_{i=1}^{n} \alpha_i k(\cdot, x_i), g = \sum_{j=1}^{m} \beta_j k(\cdot, x'_j)$. Then, $\mathcal{H}^0$ is a valid pre-Hilbert space, i.e, the pointwise closure $\mathcal{H}_k = \overline{\mathcal{H}^0}$ is a Hilbert space.*

*Proof.* We show that (3.1) indeed defines a valid inner product. First,

$$\langle f, g\rangle_{\mathcal{H}^0} = \sum_{i=1}^{n} \alpha_i g(x_i) = \sum_{j=1}^{n} \beta_j f(x'_j).$$

It is implied that that the inner product is independent of the specific representation of $f$ and $g$. The triangular inequality is easy to verify. Next, we show that $\|f\|_{\mathcal{H}^0} = 0$ if and only if $f = 0$. If there exist $x_0 \in \mathcal{X}$ such that $f(x_0) \neq 0$. Assume $f(x) = \sum_{j=1}^{m} a_j k(x_j, \cdot)$ and consider

$$0 \leq \|\lambda f + f(x_0)k(\cdot, x_0)\|^2_{\mathcal{H}^0} = \lambda^2\|f\|^2_{\mathcal{H}^0} + 2\lambda f^2(x_0) + f^2(x_0)k(x_0, x_0).$$

Taking $\lambda \to -\infty$, the RHS will be negative and this causes contradictory.

What remains is to show that the convergence of Cauchy sequence. We refer to Link for a complete proof. □

**Lemma 3.2.** *The Hilbert space defined in Theorem 3.1 satisfies the reproducing property:*

$$\langle f, k(\cdot, x)\rangle_{\mathcal{H}_k} = f(x).$$

*Proof.* For $f \in \mathcal{H}^0$, we can write $f(x) = \sum_{j=1}^{m} a_j k(\cdot, x_j)$. By definition,

$$\langle f, k(\cdot, x)\rangle_{\mathcal{H}_k} = \sum_{j=1}^{m} a_j k(x, x_j) = f(x).$$

For any $f \in \mathcal{H}_k$, let $\lim_{n\to\infty} f_n(x) = f(x)$. Then,

$$\langle f, k(\cdot, x)\rangle_{\mathcal{H}_k} = \lim_{n\to\infty} \langle f_n, k(\cdot, x)\rangle_{\mathcal{H}_k} = \lim_{n\to\infty} f_n(x) = f(x).$$

□

The reproducing property is the most important property of this Hilbert space.

**Definition 3.3** (RKHS)**.** Let $\mathcal{X}$ be an arbitrary set and $\mathcal{H}$ a Hilbert space of real-valued functions on $\mathcal{X}$. We say $\mathcal{H}$ is a reproducing kernel Hilbert space (RKHS) if there is a kernel $k : \mathcal{X} \times \mathcal{X} \mapsto \mathbb{R}$ such that

- $\forall x \in \mathcal{X}, k(\cdot, x) \in \mathcal{H}$.
- *Reproducing property:* $\forall x \in X, f \in \mathcal{H}, \langle f, k(\cdot, x)\rangle_{\mathcal{H}} = f(x)$

**Lemma 3.4.** *For a RKHS, the evaluation functional $L_x(f) = f(x)$ is continuous.*

*Proof.* For any $x \in X$ and $f \in \mathcal{H}$,

$$\sup_{\|f\|_{\mathcal{H}} \leq 1} |L_x(f)| = \sup_{\|f\|_{\mathcal{H}} \leq 1} |\langle f, k(\cdot, x)\rangle_{\mathcal{H}}| \leq \|k(\cdot, x)\|_{\mathcal{H}} < \infty.$$

□

This continuity of the evaluation functional is sometimes used as the equivalent definition of RKHS. An important implication is that the convergence in norm implies the pointwise convergence. If $\lim_{n\to\infty} \|f_n - f\|_{\mathcal{H}} = 0$, then

$$|f_n(x) - f(x)| \leq \|L_x\| \|f_n - f\|_{\mathcal{H}} \to 0 \qquad \text{as } n \to \infty.$$

**Lemma 3.5.** *For a RKHS, the reproducing kernel $k$ is unique.*

*Proof.* For any two kernels $k_1, k_2$,

$$\langle f, k_1(\cdot, x) - k_2(\cdot, x)\rangle_{\mathcal{H}} = f(x) - f(x) = 0, \forall x \in X, \forall\, f \in \mathcal{H}.$$

Taking $f = k_1(\cdot, x) - k_2(\cdot, x)$, we have $\|k_1(\cdot, x) - k_2(\cdot, x)\|_{\mathcal{H}}^2 = 0, \forall x \in X$. Hence, $k_1 = k_2$. □

**Theorem 3.6.** *For any kernel $k$, there is a unique RKHS, for which $k$ is the reproducing kernel.*

*Proof.* First, by Moore-Aronsajn theorem, there exists a RKHS with $k$ being the reproducing kernel. Assume $\mathcal{H}_1$ and $\mathcal{H}_2$ be two RKHSs with $k$ being the reproducing kernel. First, by definition, $k(\cdot, x) \in \mathcal{H}_1$ for any $x \in \mathcal{X}$. Hence, $\mathcal{H}^0 \subset \mathcal{H}_1$. Moreover, $\mathcal{H}^0$ is dense in $\mathcal{H}_1$ since if there exists $f \in \mathcal{H}$ such that $f \perp \mathcal{H}^0$, we must have

$$\langle f, k(\cdot, x)\rangle_{\mathcal{H}_1} = f(x) = 0 \qquad \forall x \in \mathcal{X}.$$

For $f = \sum_{j=1}^m a_j k(\cdot, x_j)$,

$$\begin{aligned}\|f\|_{\mathcal{H}_1}^2 &= \langle \sum_i^n a_i k(\cdot, x_i), \sum_{j=1}^m a_j k(\cdot, x_j)\rangle_{\mathcal{H}_1} = \sum_{i,j=1}^n a_i a_j \langle k(\cdot, x_i), k(\cdot, x_j)\rangle_{\mathcal{H}_1} \\ &\overset{(i)}{=} \sum_{i,j=1}^n a_i a_j k(x_i, x_j) = \|f\|_{\mathcal{H}^0}^2.\end{aligned}$$

where $(i)$ follows from the reproducing property. Hence, $\|f\|_{\mathcal{H}_1} = \|f\|_{\mathcal{H}^0}$ for $f \in \mathcal{H}_0$. By the same argument, the same results hold for $\mathcal{H}_2$. For any $f \in \mathcal{H}_1$, there must exits $(f_n) \subset \mathcal{H}^0$ such that $f(x) = \lim_{n\to\infty} f_n(x)$. This implies that $f \in \mathcal{H}_2$. Similarly, $\mathcal{H}_1$ and $\mathcal{H}_2$ contains the same functions. What remains is to check that the two norms coincide, which results from

$$\|f\|_{\mathcal{H}_1} = \lim_{n\to\infty} \|f_n\|_{\mathcal{H}_1} = \lim_{n\to\infty} \|f_n\|_{\mathcal{H}^0} = \lim_{n\to\infty} \|f_n\|_{\mathcal{H}_2} = \|f\|_{\mathcal{H}_2}.$$

□

**Theorem 3.7.** *A Hilbert space of functions $\mathcal{H} \subset \mathbb{R}^{\mathcal{X}}$ is a RKHS if and only if the evaluation functional is continuous.*

*Proof.* If $L_x$ is continuous, by Riesz representation theorem, there exist $K_x \in \mathcal{H}$ such that

$$L_x(f) = \langle K_x, f \rangle_{\mathcal{H}}.$$

Define the kernel:

$$k(x, x') = \langle K_x, K_{x'} \rangle_{\mathcal{H}} = K_{x'}(x) = K_x(x'),$$

for which

$$\langle f, k(\cdot, x) \rangle_{\mathcal{H}} = \langle f, K_x \rangle = f(x), \quad \forall f \in \mathcal{H}.$$

This means $k(\cdot, \cdot)$ is a reproducing kernel of $\mathcal{H}$. □

# 4 A generalization analysis of kernel ridge regression

We first provide the upper bound of the Rademacher complexity.

**Proposition 4.1.** *For any kernel $k$, let $\mathcal{H}_k$ the corresponding RKHS. Let $\mathcal{H}_k^Q = \{f \in \mathcal{H}_k : \|f\|_{\mathcal{H}_k} \leq Q\}$. Then, we have*

$$\widehat{\mathrm{Rad}}_n(\mathcal{H}_k^Q) \leq Q \frac{\sqrt{\sum_{i=1}^n k(x_i, x_i)}}{n}.$$

*Proof.*

$$\begin{aligned}
n\widehat{\mathrm{Rad}}_n(\mathcal{H}_k^Q) &= \mathbb{E}_\xi[\sup_{\|f\|_{\mathcal{H}_k} \leq Q} \sum_{i=1}^n \xi_i f(x_i)] = \mathbb{E}_\xi[\sup_{\|f\|_{\mathcal{H}_k} \leq Q} \sum_{i=1}^n \xi_i \langle f, k(\cdot, x_i) \rangle_{\mathcal{H}_k}] \text{(reproducing property)} \\
&= \mathbb{E}_\xi[\sup_{\|f\|_{\mathcal{H}_k} \leq Q} \langle f, \sum_{i=1}^n \xi_i k(\cdot, x_i) \rangle_{\mathcal{H}}] \leq Q\, \mathbb{E}_\xi[\| \sum_{i=1}^n \xi_i k(\cdot, x_i) \|_{\mathcal{H}_k}] \\
&= Q\, \mathbb{E}_\xi \sqrt{\sum_{i,j=1}^n \xi_i \xi_j k(x_i, x_j)} \leq Q \sqrt{\mathbb{E}_\xi[\sum_{i,j=1}^n \xi_i \xi_j k(x_i, x_j)]} \qquad \text{(Jensen inequality)} \\
&= Q \sqrt{\sum_{i=1}^n k(x_i, x_i)} \qquad (\mathbb{E}[\xi_i \xi_j] = 0, \forall\, i \neq j).
\end{aligned}$$

□

Given data $\{(x_i, f^*(x_i))\}_{i=1}^n$, consider the kernel ridge regression (KRR) estimator

$$\hat{f}_n = \operatorname*{argmin}_{f \in \mathcal{H}_k} \hat{\mathcal{R}}_n(f) + \lambda \|f\|_{\mathcal{H}_k}. \tag{4.1}$$

**Theorem 4.2** (A priori estimate)**.** *Assume that $\ell(\cdot, y)$ is $L$-Lipschitz and bounded by $B$, and $\sup_{x \in \mathcal{X}} k(x, x) \leq 1$. Then, for any $\delta \in (0, 1)$, with probability $1 - \delta$ over the choice of training set, we have*

$$\mathcal{R}(\hat{f}_n) \lesssim \lambda \|f^*\|_{\mathcal{H}_k} + \frac{L \|f^*\|_{\mathcal{H}_k}}{\sqrt{n}} + B \sqrt{\frac{\log(1/\delta)}{n}}.$$

*Proof.* (1) Let $Q = \|f^*\|_{\mathcal{H}_k}$. By the definition of $\hat{f}_n$,

$$\hat{\mathcal{R}}_n(\hat{f}_n) + \lambda\|\hat{f}_n\|_{\mathcal{H}_k} \le \hat{\mathcal{R}}_n(f^*) + \lambda\|f^*\|_{\mathcal{H}_k} = \lambda\|f^*\|_{\mathcal{H}_k} = \lambda Q,$$

which yields

$$\|\hat{f}_n\|_{\mathcal{H}_k} \le Q, \qquad \hat{\mathcal{R}}_n(\hat{f}_n) \le \lambda Q.$$

(2) Let $\mathcal{F}_Q = \{\ell(h(x), h^*(x)) : h \in \mathcal{H}_k^Q\}$. By the contraction lemma, we have

$$\hat{\mathcal{R}}_n(\mathcal{F}_Q) \le L\hat{\mathcal{R}}_n(\mathcal{H}_k^Q).$$

Using the Rademacher complexity-based generalization bound, we have

$$\begin{aligned}|\hat{\mathcal{R}}_n(\hat{f}_n) - \mathcal{R}(\hat{f}_n)| &\le \sup_{\|f\|_{\mathcal{H}} \le Q} |\hat{\mathcal{R}}_n(f) - \mathcal{R}(f)| \lesssim \hat{\mathcal{R}}_n(\mathcal{F}_Q) + B\sqrt{\frac{\log(4/\delta)}{n}} \\ &\lesssim L\hat{\mathcal{R}}_n(\mathcal{H}_k^Q) + B\sqrt{\frac{\log(4/\delta)}{n}} \le \frac{LQ}{\sqrt{n}} + B\sqrt{\frac{\log(1/\delta)}{n}} \quad (\text{use } \sup_{x\in X} k(x,x) \le 1).\end{aligned}$$

(3) $\mathcal{R}(\hat{f}_n) \le \hat{\mathcal{R}}_n(\hat{f}_n) + |\hat{\mathcal{R}}_n(\hat{f}_n) - \mathcal{R}(\hat{f}_n)| \le \lambda Q + (LQ + B\sqrt{\log(4/\delta)})/\sqrt{n}$.

□

The preceding estimate is a priori, since it depends on the norm of $f^*$ instead of that of $\hat{f}_n$. Taking $\lambda = O(1/\sqrt{n})$, we have that $\mathcal{R}(\hat{f}_n) = O(1/\sqrt{n})$, which does not suffer from the curse of dimensionality. This means that the functions in the RKHS can be efficiently learned by the KRR.

- Similar results hold for any regularizations of the form $r(\|f\|_{\mathcal{H}_k})$, where $r : [0,\infty) \to [0,\infty)$ is strictly increasing.
- Note that Theorem 4.2 holds as long as $\lambda > 0$ and one can even take $\lambda \to 0^+$, which may seem strange at the first glance. This is due to that there is no label noise. In fact, the optimal $\lambda$ depends on the level of noise as shown in the following theorem.

Consider the estimator

$$\hat{f}_n = \operatorname*{argmin}_{f\in\mathcal{H}_k} \frac{1}{n}\sum_{i=1}^n (T \circ f(x_i) - y_i)^2 + \lambda \|f\|_{\mathcal{H}_k}, \tag{4.2}$$

where $T(t) = \min(\max(t,-1),1)$. We make the following non-essential technical assumptions.

- $\sup_x |f^*(x)| \le 1$ and $\sup_x k(x,x) \le 1$.
- $y_i = f^*(x_i) + \xi_i$. $\{\xi\}_i$ are i.i.d. random noises with $|\xi_i| \le \sigma$. Assume that $\sigma \le 1$.

**Theorem 4.3.** *Under the preceding assumptions and taking $\lambda = \frac{\sigma}{\sqrt{n}}$, for any $\delta \in (0,1)$, we have*

$$\|\hat{f}_n - f^*\|^2_{L^2(\mathbb{P}_x)} \lesssim \frac{\|f^*\|_{\mathcal{H}_k} + \sqrt{\log(2/\delta)}}{\sqrt{n}}.$$

*Proof.* Let $Q = \|f^*\|_{\mathcal{H}_k}$. By definition,

$$\hat{\mathcal{R}}_n(\hat{f}_n) + \lambda\|\hat{f}_n\|_{\mathcal{H}_k} \le \hat{\mathcal{R}}_n(f^*) + \lambda\|f^*\|_{\mathcal{H}_k} = \hat{\mathcal{R}}_n(f^*) + \lambda Q.$$

which yields

$$\hat{\mathcal{R}}_n(\hat{f}_n) \le \hat{\mathcal{R}}_n(f^*) + \lambda Q$$
$$\|\hat{f}_n\|_{\mathcal{H}_k} \le Q + \frac{\hat{\mathcal{R}}_n(f^*)}{\lambda}.$$

Notice that $|y_i| \le 1+\sigma$, hence $\phi_i(t) := (t-y_i)^2$ is $2(2+\sigma)$-Lipschitz continuous. Using the contraction lemma and Rademacher complexity-based bound, for any $\delta_1 \in (0,1)$, we have with probability $1-\delta_1$,

$$\begin{aligned}
\mathcal{R}(\hat{f}_n) &\le \hat{\mathcal{R}}_n(\hat{f}_n) + \sup_{\|f\|_{\mathcal{H}_k} \le Q + \frac{\hat{\mathcal{R}}_n(f^*)}{\lambda}} |\mathcal{R}(f) - \hat{\mathcal{R}}_n(f)| \\
&\lesssim \hat{\mathcal{R}}_n(f^*) + \lambda Q + 2(2+\sigma)\,\mathrm{Rad}_n(\mathcal{F}_{Q+\frac{\hat{\mathcal{R}}_n(f^*)}{\lambda}}) + (1+\sigma)\sqrt{\frac{\log(2/\delta_1)}{n}} \\
&\lesssim \hat{\mathcal{R}}_n(f^*) + \lambda Q + \frac{Q + \frac{\hat{\mathcal{R}}_n(f^*)}{\lambda}}{\sqrt{n}} + \sqrt{\frac{\log(2/\delta_1)}{n}}. \qquad (4.3)
\end{aligned}$$

Notice that

$$\hat{\mathcal{R}}_n(\hat{f}_n) = \mathbb{E}[(\hat{f}_n(x) - f^*(x) - \xi)] = \|\hat{f}_n - f^*\|^2_{L^2(\mathbb{P}_x)} + \mathbb{E}[\xi^2]$$
$$\hat{\mathcal{R}}_n(f^*) = \frac{1}{n}\sum_{i=1}^n \xi_i^2.$$

By Hoeffding's inequality, for any $\delta_2 \in (0,1)$, with probability at least $1-\delta_2$,

$$\frac{1}{n}\sum_{i=1}^n \xi_i^2 - \mathbb{E}[\xi^2] \le \sigma\sqrt{\frac{\log(1/\delta_2)}{n}}.$$

Plugging it into (4.3),

$$\begin{aligned}
\|\hat{f}_n - f^*\|^2_{L^2} &\lesssim \frac{1}{n}\sum_{i=1}^n \xi_i^2 - \mathbb{E}[\xi^2] + \lambda Q + \frac{\frac{1}{n}\sum_{i=1}^n \xi_i^2}{\lambda\sqrt{n}} + \frac{Q}{\sqrt{n}} + \sqrt{\frac{\log(1/\delta_1)}{n}} \\
&\lesssim \lambda Q + \frac{\sigma\sqrt{\log(1/\delta_2)}}{\lambda n} + \frac{Q}{\sqrt{n}} + \sqrt{\frac{\log(1/\delta_1)}{n}}.
\end{aligned}$$

Taking $\lambda = \frac{\sigma}{\sqrt{n}}$ and $\delta_2 = \delta_1 = \delta/2$, we complete the proof. □

**Tightness.** Note that the preceding bounds are not tight for the square loss: $\ell(y_1, y_2) = (y_1 - y_2)^2$. When applying the contraction lemma, we use the worst-case Lipschitz norm $\mathrm{Lip}(t^2/2) \le 1$ for $t \in [0,1]$. However, at the estimator, we should have $\varepsilon(x) = \hat{f}(x) - f^*(x) \ll 1$. Therefore, we should use the "local" Lipschitz norm to bound the Rademacher complexity. This will in turn gives rise to a fast rate. Usually, the fast rate close to $O(1/n)$ and this approach is called "local Rademacher complexity". Please refer to [Bartlett et al., 2005] for more details.

## References

[Bartlett et al., 2005] Bartlett, P. L., Bousquet, O., Mendelson, S., et al. (2005). Local Rademacher complexities. *The Annals of Statistics*, 33(4):1497–1537.